

data-wrangling

October 9, 2019

```
<a href="https://cocl.us/corsera_da0101en_notebook_top">  
    
</a>
```

Data Analysis with Python

Data Wrangling

Welcome!

By the end of this notebook, you will have learned the basics of Data Wrangling!

Table of content

Identify and handle missing values

Identify missing values

Deal with missing values

Correct data format

Data standardization

Data Normalization (centering/scaling)

Binning

Indicator variable

Estimated Time Needed: 30 min

What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

What is the fuel consumption (L/100k) rate for the diesel car?

Import data

You can find the “Automobile Data Set” from the following link:
<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will
be using this data set throughout this course.

Import pandas

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
```

Reading the data set from the URL and adding the related headers.

URL of the dataset

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[2]: filename = "https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/
↳CognitiveClass/DA0101EN/auto.csv"
```

Python list headers containing name of headers

```
[3]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
↳"num-of-doors", "body-style",
"drive-wheels", "engine-location", "wheel-base",
↳"length", "width", "height", "curb-weight", "engine-type",
"num-of-cylinders",
↳"engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
"peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter “names” equal to the Python list “headers”.

```
[4]: df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
[5]: # To see what the data set looks like, we'll use the head() method.
df.head()
```

```
[5]:
```

| | symboling | normalized-losses | make | fuel-type | aspiration | num-of-doors | \ |
|---|-----------|-------------------|-------------|-----------|------------|--------------|---|
| 0 | 3 | ? | alfa-romero | gas | std | two | |
| 1 | 3 | ? | alfa-romero | gas | std | two | |
| 2 | 1 | ? | alfa-romero | gas | std | two | |
| 3 | 2 | 164 | audi | gas | std | four | |
| 4 | 2 | 164 | audi | gas | std | four | |

| | body-style | drive-wheels | engine-location | wheel-base | ... | engine-size | \ |
|---|-------------|--------------|-----------------|------------|-----|-------------|---|
| 0 | convertible | rwd | front | 88.6 | ... | 130 | |
| 1 | convertible | rwd | front | 88.6 | ... | 130 | |
| 2 | hatchback | rwd | front | 94.5 | ... | 152 | |
| 3 | sedan | fwd | front | 99.8 | ... | 109 | |
| 4 | sedan | 4wd | front | 99.4 | ... | 136 | |

| | fuel-system | bore | stroke | compression-ratio | horsepower | peak-rpm | city-mpg | \ |
|---|-------------|------|--------|-------------------|------------|----------|----------|---|
| 0 | mpfi | 3.47 | 2.68 | 9.0 | 111 | 5000 | 21 | |
| 1 | mpfi | 3.47 | 2.68 | 9.0 | 111 | 5000 | 21 | |

| | | | | | | | |
|---|------|------|------|------|-----|------|----|
| 2 | mpfi | 2.68 | 3.47 | 9.0 | 154 | 5000 | 19 |
| 3 | mpfi | 3.19 | 3.40 | 10.0 | 102 | 5500 | 24 |
| 4 | mpfi | 3.19 | 3.40 | 8.0 | 115 | 5500 | 18 |

| | highway-mpg | price |
|---|-------------|-------|
| 0 | 27 | 13495 |
| 1 | 27 | 16500 |
| 2 | 26 | 16500 |
| 3 | 30 | 13950 |
| 4 | 22 | 17450 |

[5 rows x 26 columns]

As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data:

identify missing data

deal with missing data

correct data format

Identify and handle missing values

Identify missing values

Convert “?” to NaN

In the car dataset, missing data comes with the question mark “?”. We replace “?” with NaN (Not a Number), which is Python’s default missing value marker, for reasons of computational speed and convenience. Here we use the function:

to replace A by B

```
[6]: import numpy as np

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)
```

```
[6]:   symboling  normalized-losses   make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero    gas      std         two
1         3             NaN  alfa-romero    gas      std         two
2         1             NaN  alfa-romero    gas      std         two
3         2           164      audi      gas      std         four
4         2           164      audi      gas      std         four
```

| | body-style | drive-wheels | engine-location | wheel-base | ... | engine-size | \ |
|---|-------------|--------------|-----------------|------------|-----|-------------|---|
| 0 | convertible | rwd | front | 88.6 | ... | 130 | |
| 1 | convertible | rwd | front | 88.6 | ... | 130 | |
| 2 | hatchback | rwd | front | 94.5 | ... | 152 | |
| 3 | sedan | fwd | front | 99.8 | ... | 109 | |
| 4 | sedan | 4wd | front | 99.4 | ... | 136 | |

| | fuel-system | bore | stroke | compression-ratio | horsepower | peak-rpm | city-mpg | \ |
|---|-------------|------|--------|-------------------|------------|----------|----------|---|
| 0 | mpfi | 3.47 | 2.68 | 9.0 | 111 | 5000 | 21 | |
| 1 | mpfi | 3.47 | 2.68 | 9.0 | 111 | 5000 | 21 | |
| 2 | mpfi | 2.68 | 3.47 | 9.0 | 154 | 5000 | 19 | |
| 3 | mpfi | 3.19 | 3.40 | 10.0 | 102 | 5500 | 24 | |
| 4 | mpfi | 3.19 | 3.40 | 8.0 | 115 | 5500 | 18 | |

| | highway-mpg | price |
|---|-------------|-------|
| 0 | 27 | 13495 |
| 1 | 27 | 16500 |
| 2 | 26 | 16500 |
| 3 | 30 | 13950 |
| 4 | 22 | 17450 |

[5 rows x 26 columns]

identify__missing__values

Evaluating for Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. There are two methods to detect missing data:

```
.isnull()
.notnull()
```

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
[7]: missing_data = df.isnull()
missing_data.head(5)
```

| | symboling | normalized-losses | make | fuel-type | aspiration | num-of-doors | \ |
|---|-----------|-------------------|-------|-----------|------------|--------------|---|
| 0 | False | True | False | False | False | False | |
| 1 | False | True | False | False | False | False | |
| 2 | False | True | False | False | False | False | |
| 3 | False | False | False | False | False | False | |
| 4 | False | False | False | False | False | False | |

| | body-style | drive-wheels | engine-location | wheel-base | ... | engine-size | \ |
|---|------------|--------------|-----------------|------------|-----|-------------|---|
| 0 | False | False | False | False | ... | False | |
| 1 | False | False | False | False | ... | False | |

| | | | | | | |
|---|-------|-------|-------|-------|-------|-------|
| 2 | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False |

| | fuel-system | bore | stroke | compression-ratio | horsepower | peak-rpm | \ |
|---|-------------|-------|--------|-------------------|------------|----------|---|
| 0 | False | False | False | False | False | False | |
| 1 | False | False | False | False | False | False | |
| 2 | False | False | False | False | False | False | |
| 3 | False | False | False | False | False | False | |
| 4 | False | False | False | False | False | False | |

| | city-mpg | highway-mpg | price |
|---|----------|-------------|-------|
| 0 | False | False | False |
| 1 | False | False | False |
| 2 | False | False | False |
| 3 | False | False | False |
| 4 | False | False | False |

[5 rows x 26 columns]

“True” stands for missing value, while “False” stands for not missing value.

Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, “True” represents a missing value, “False” means the value is present in the dataset. In the body of the for loop the method “.value_counts()” counts the number of “True” values.

```
[8]: for column in missing_data.columns.values.tolist():
      print(column)
      print(missing_data[column].value_counts())
      print("")
```

```
symboling
False      205
Name: symboling, dtype: int64
```

```
normalized-losses
False      164
True        41
Name: normalized-losses, dtype: int64
```

```
make
False      205
Name: make, dtype: int64
```

```
fuel-type
False      205
```

Name: fuel-type, dtype: int64

aspiration
False 205
Name: aspiration, dtype: int64

num-of-doors
False 203
True 2
Name: num-of-doors, dtype: int64

body-style
False 205
Name: body-style, dtype: int64

drive-wheels
False 205
Name: drive-wheels, dtype: int64

engine-location
False 205
Name: engine-location, dtype: int64

wheel-base
False 205
Name: wheel-base, dtype: int64

length
False 205
Name: length, dtype: int64

width
False 205
Name: width, dtype: int64

height
False 205
Name: height, dtype: int64

curb-weight
False 205
Name: curb-weight, dtype: int64

engine-type
False 205
Name: engine-type, dtype: int64

num-of-cylinders

False 205
Name: num-of-cylinders, dtype: int64

engine-size
False 205
Name: engine-size, dtype: int64

fuel-system
False 205
Name: fuel-system, dtype: int64

bore
False 201
True 4
Name: bore, dtype: int64

stroke
False 201
True 4
Name: stroke, dtype: int64

compression-ratio
False 205
Name: compression-ratio, dtype: int64

horsepower
False 203
True 2
Name: horsepower, dtype: int64

peak-rpm
False 203
True 2
Name: peak-rpm, dtype: int64

city-mpg
False 205
Name: city-mpg, dtype: int64

highway-mpg
False 205
Name: highway-mpg, dtype: int64

price
False 201
True 4
Name: price, dtype: int64

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

“normalized-losses”: 41 missing data

“num-of-doors”: 2 missing data

“bore”: 4 missing data

“stroke” : 4 missing data

“horsepower”: 2 missing data

“peak-rpm”: 2 missing data

“price”: 4 missing data

Deal with missing data

How to deal with missing data?

drop data a. drop the whole row b. drop the whole column

replace data a. replace it by mean b. replace it by frequency c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

“normalized-losses”: 41 missing data, replace them with mean

“stroke”: 4 missing data, replace them with mean

“bore”: 4 missing data, replace them with mean

“horsepower”: 2 missing data, replace them with mean

“peak-rpm”: 2 missing data, replace them with mean

Replace by frequency:

“num-of-doors”: 2 missing data, replace them with “four”.

Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

Drop the whole row:

“price”: 4 missing data, simply delete the whole row

Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

Calculate the average of the column

```
[9]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```


Average of normalized-losses: 122.0

Replace “NaN” by mean value in “normalized-losses” column

```
[10]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

Calculate the mean value for ‘bore’ column

```
[11]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

Replace NaN by mean value

```
[12]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

Question #1:

According to the example above, replace NaN in “stroke” column by mean.

```
[13]: # Write your code below and press Shift+Enter to execute

avg_stroke=df["stroke"].astype('float').mean(axis=0)
print("Average Stroke:", avg_stroke)
```

Average Stroke: 3.255422885572139

Double-click here for the solution.

Calculate the mean value for the ‘horsepower’ column:

```
[14]: avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
      print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

Replace “NaN” by mean value:

```
[15]: df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

Calculate the mean value for ‘peak-rpm’ column:

```
[16]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
      print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

Replace NaN by mean value:

```
[17]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the “value_counts()” method:

```
[18]: df['num-of-doors'].value_counts()
```

```
[18]: four      114
      two       89
      Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the “idxmax()” method to calculate for us the most common type automatically:

```
[19]: df['num-of-doors'].value_counts().idxmax()
```

```
[19]: 'four'
```

The replacement procedure is very similar to what we have seen previously

```
[20]: #replace the missing 'num-of-doors' values by the most frequent
      df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let’s drop all rows that do not have price data:

```
[21]: # simply drop whole row with NaN in "price" column
      df.dropna(subset=["price"], axis=0, inplace=True)

      # reset index, because we dropped two rows
      df.reset_index(drop=True, inplace=True)
```

```
[22]: df.head()
```

```
[22]:      symboling  normalized-losses      make fuel-type aspiration num-of-doors  \
0           3           122  alfa-romero    gas      std         two
1           3           122  alfa-romero    gas      std         two
2           1           122  alfa-romero    gas      std         two
3           2           164      audi     gas      std         four
4           2           164      audi     gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front        88.6  ...        130
1  convertible         rwd         front        88.6  ...        130
2   hatchback         rwd         front        94.5  ...        152
3      sedan         fwd         front        99.8  ...        109
4      sedan         4wd         front        99.4  ...        136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi    3.47    2.68           9.0         111      5000      21
1      mpfi    3.47    2.68           9.0         111      5000      21
2      mpfi    2.68    3.47           9.0         154      5000      19
3      mpfi    3.19    3.40          10.0         102      5500      24
4      mpfi    3.19    3.40           8.0         115      5500      18
```

| | highway-mpg | price |
|---|-------------|-------|
| 0 | 27 | 13495 |
| 1 | 27 | 16500 |
| 2 | 26 | 16500 |
| 3 | 30 | 13950 |
| 4 | 22 | 17450 |

[5 rows x 26 columns]

Good! Now, we obtain the dataset with no missing values.

Correct data format

We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

.dtype() to check the data type

.astype() to change the data type

Lets list the data types for each column

```
[23]: df.dtypes
```

```
[23]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                 float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   object
stroke                 object
compression-ratio      float64
horsepower             object
```

```

peak-rpm          object
city-mpg           int64
highway-mpg        int64
price             object
dtype: object

```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

Convert data types to proper format

```

[24]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
      df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
      df[["price"]] = df[["price"]].astype("float")
      df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")

```

Let us list the columns after the conversion

```

[25]: df.dtypes

```

```

[25]: symboling          int64
      normalized-losses  int64
      make              object
      fuel-type         object
      aspiration        object
      num-of-doors      object
      body-style        object
      drive-wheels      object
      engine-location   object
      wheel-base        float64
      length            float64
      width             float64
      height            float64
      curb-weight       int64
      engine-type       object
      num-of-cylinders  object
      engine-size       int64
      fuel-system       object
      bore              float64
      stroke            float64
      compression-ratio float64
      horsepower       object
      peak-rpm         float64
      city-mpg         int64
      highway-mpg      int64

```

```
price          float64
dtype: object
```

Wonderful!

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns “city-mpg” and “highway-mpg” are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard

We will need to apply data transformation to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

```
[26]: df.head()
```

```
[26]:   symboling  normalized-losses      make fuel-type aspiration \
0         3         122  alfa-romero    gas      std
1         3         122  alfa-romero    gas      std
2         1         122  alfa-romero    gas      std
3         2         164      audi    gas      std
4         2         164      audi    gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...

   engine-size  fuel-system  bore  stroke  compression-ratio  horsepower  \
0         130      mpfi  3.47   2.68           9.0          111
1         130      mpfi  3.47   2.68           9.0          111
```

| | | | | | | |
|---|-----|------|------|------|------|-----|
| 2 | 152 | mpfi | 2.68 | 3.47 | 9.0 | 154 |
| 3 | 109 | mpfi | 3.19 | 3.40 | 10.0 | 102 |
| 4 | 136 | mpfi | 3.19 | 3.40 | 8.0 | 115 |

| | peak-rpm | city-mpg | highway-mpg | price |
|---|----------|----------|-------------|---------|
| 0 | 5000.0 | 21 | 27 | 13495.0 |
| 1 | 5000.0 | 21 | 27 | 16500.0 |
| 2 | 5000.0 | 19 | 26 | 16500.0 |
| 3 | 5500.0 | 24 | 30 | 13950.0 |
| 4 | 5500.0 | 18 | 22 | 17450.0 |

[5 rows x 26 columns]

```
[27]: # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

```
[27]:      symboling  normalized-losses      make fuel-type aspiration \
0          3           122  alfa-romero      gas      std
1          3           122  alfa-romero      gas      std
2          1           122  alfa-romero      gas      std
3          2           164      audi      gas      std
4          2           164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0          mpfi  3.47   2.68              9.0          111   5000.0      21
1          mpfi  3.47   2.68              9.0          111   5000.0      21
2          mpfi  2.68   3.47              9.0          154   5000.0      19
3          mpfi  3.19   3.40             10.0          102   5500.0      24
4          mpfi  3.19   3.40              8.0          115   5500.0      18

      highway-mpg  price  city-L/100km
0          27  13495.0    11.190476
1          27  16500.0    11.190476
2          26  16500.0    12.368421
3          30  13950.0     9.791667
4          22  17450.0    13.055556
```

[5 rows x 27 columns]

Question #2:

According to the example above, transform mpg to L/100km in the column of “highway-mpg”, and change the name of column to “highway-L/100km”.

```
[29]: # Write your code below and press Shift+Enter to execute

# transform mpg to L/100km by mathematical operation (235 divided by mpg)
df["highway-mpg"] = 235/df["highway-mpg"]

# rename column name from "highway-mpg" to "highway-L/100km"
df.rename(columns={"highway-mpg": "highway-L/100km"}, inplace=True)

# check your transformed data
df.head()
```

```
[29]:   symboling  normalized-losses      make fuel-type aspiration \
0         3             122  alfa-romero      gas      std
1         3             122  alfa-romero      gas      std
2         1             122  alfa-romero      gas      std
3         2             164      audi      gas      std
4         2             164      audi      gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible         rwd         front      88.6  ...
1         two  convertible         rwd         front      88.6  ...
2         two   hatchback         rwd         front      94.5  ...
3         four      sedan         fwd         front      99.8  ...
4         four      sedan         4wd         front      99.4  ...

   fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg \
0         mpfi  3.47   2.68              9.0         111    5000.0      21
1         mpfi  3.47   2.68              9.0         111    5000.0      21
2         mpfi  2.68   3.47              9.0         154    5000.0      19
3         mpfi  3.19   3.40             10.0         102    5500.0      24
4         mpfi  3.19   3.40              8.0         115    5500.0      18

   highway-mpg  price  city-L/100km
0         27.0  13495.0    11.190476
1         27.0  16500.0    11.190476
2         26.0  16500.0    12.368421
3         30.0  13950.0     9.791667
4         22.0  17450.0    13.055556
```

[5 rows x 27 columns]

[Double-click here for the solution.](#)

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height"

Target: would like to Normalize those variables so their value ranges from 0 to 1.

Approach: replace original value by (original value)/(maximum value)

```
[30]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Question #3:

According to the example above, normalize the column "height".

```
[36]: # Write your code below and press Shift+Enter to execute

df["height"] = df["height"]/df["height"].max()
df[["length","width","height"]].head()
```

```
[36]:      length      width      height
0  0.811148  0.890278  0.816054
1  0.811148  0.890278  0.816054
2  0.822681  0.909722  0.876254
3  0.848630  0.919444  0.908027
4  0.848630  0.922222  0.908027
```

[Double-click here for the solution.](#)

Here we can see, we've normalized "length", "width" and "height" in the range of [0,1].

Binning

Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical 'bins', for grouped analysis.

Example:

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium

horsepower, and little horsepower (3 types)? Can we rearrange them into three ‘bins’ to simplify analysis?

We will use the Pandas method ‘cut’ to segment the ‘horsepower’ column into 3 bins

Example of Binning Data In Pandas

Convert data to correct format

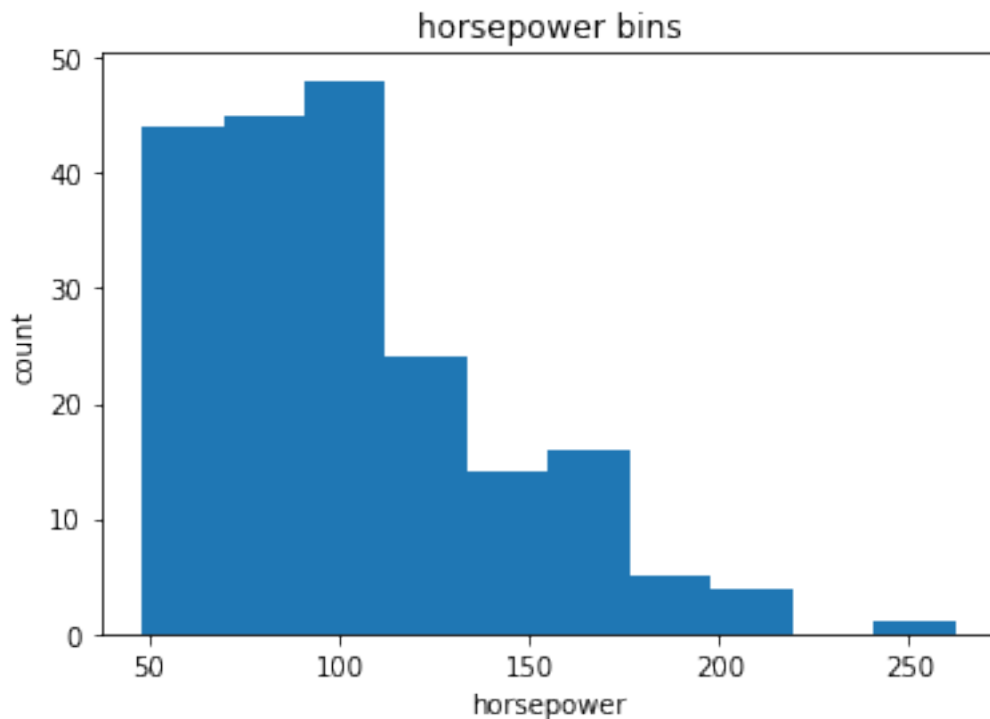
```
[38]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Lets plot the histogram of horspower, to see what the distribution of horsepower looks like.

```
[39]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[39]: Text(0.5, 1.0, 'horsepower bins')
```



We would like 3 bins of equal size bandwidth so we use numpy’s linspace(start_value, end_value,

numbers_generated function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
[40]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[40]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
[41]: group_names = ['Low', 'Medium', 'High']
```

We apply the function “cut” to determine what each value of “df[‘horsepower’]” belongs to.

```
[42]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
    ↪ include_lowest=True )
      df[['horsepower', 'horsepower-binned']].head(20)
```

```
[42]:
```

| | horsepower | horsepower-binned |
|----|------------|-------------------|
| 0 | 111 | Low |
| 1 | 111 | Low |
| 2 | 154 | Medium |
| 3 | 102 | Low |
| 4 | 115 | Low |
| 5 | 110 | Low |
| 6 | 110 | Low |
| 7 | 110 | Low |
| 8 | 140 | Medium |
| 9 | 101 | Low |
| 10 | 101 | Low |
| 11 | 121 | Medium |
| 12 | 121 | Medium |
| 13 | 121 | Medium |
| 14 | 182 | Medium |
| 15 | 182 | Medium |
| 16 | 182 | Medium |
| 17 | 48 | Low |
| 18 | 70 | Low |
| 19 | 70 | Low |

Lets see the number of vehicles in each bin.

```
[43]: df["horsepower-binned"].value_counts()
```

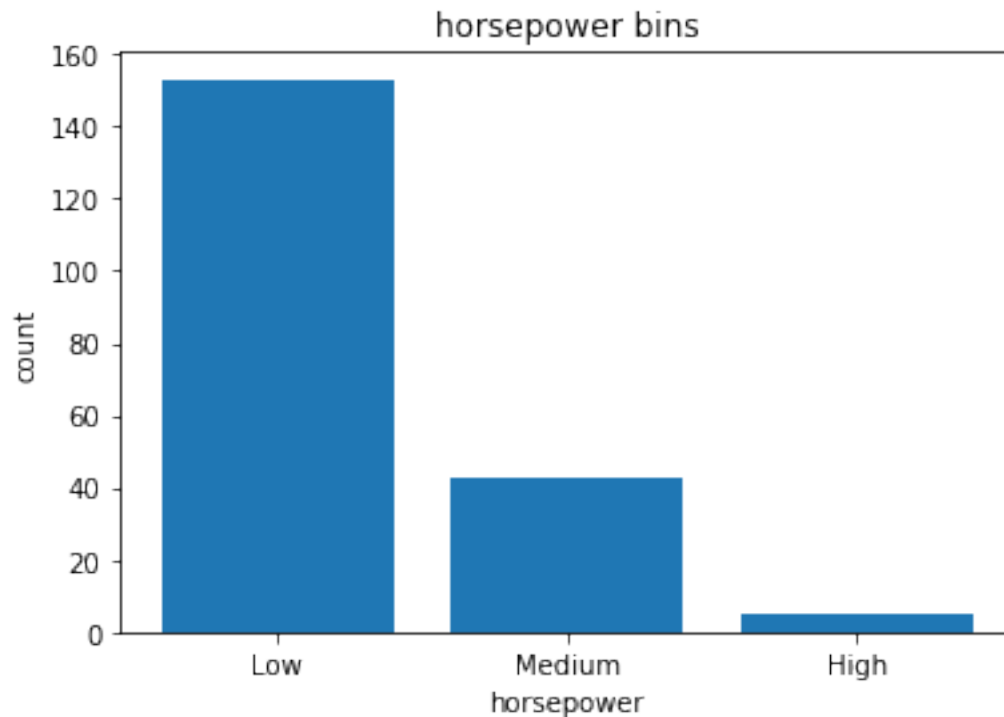
```
[43]: Low      153
      Medium   43
      High     5
      Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
[44]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[44]: Text(0.5, 1.0, 'horsepower bins')
```



Check the dataframe above carefully, you will find the last column provides the bins for “horsepower” with 3 categories (“Low”, “Medium” and “High”).

We successfully narrow the intervals from 57 to 3!

Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

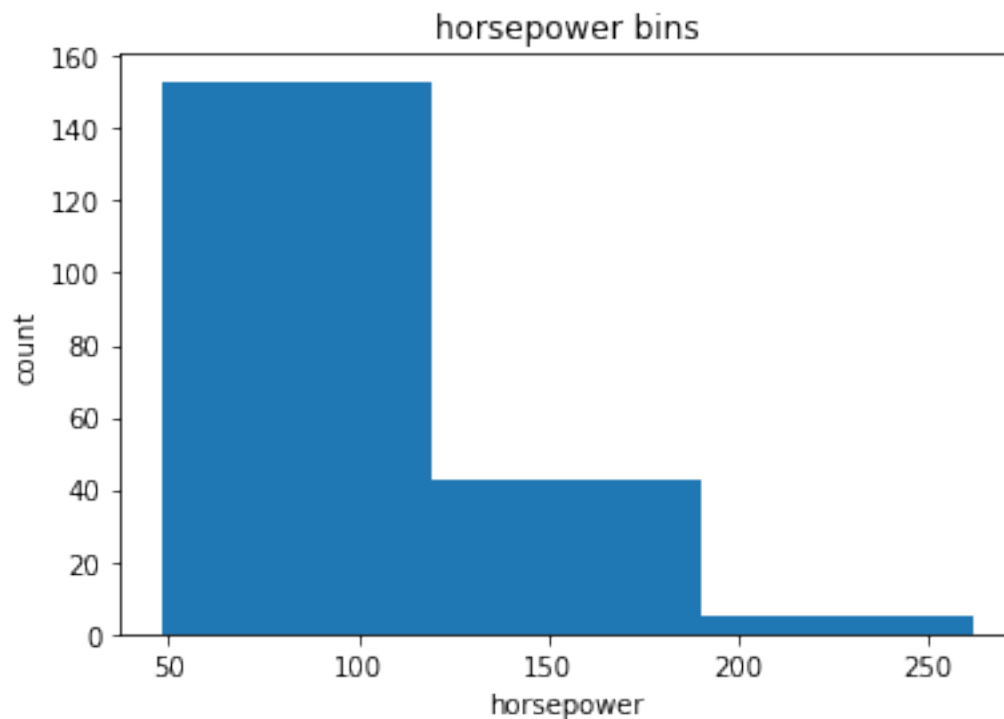
```
[45]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot

a = (0,1,2)

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[45]: Text(0.5, 1.0, 'horsepower bins')
```



The plot above shows the binning result for attribute “horsepower”.

Indicator variable (or dummy variable)

What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called ‘dummies’ because the numbers themselves don’t have inherent meaning.

Why we use indicator variables?

So we can use categorical variables for regression analysis in the later modules.

Example

We see the column “fuel-type” has two unique values, “gas” or “diesel”. Regression doesn’t understand words, only numbers. To use this attribute in regression analysis, we convert “fuel-type” into indicator variables.

We will use the panda’s method ‘get_dummies’ to assign numerical values to different categories of fuel type.

```
[52]: df.columns
```

```
[52]: Index(['symboling', 'normalized-losses', 'make', 'aspiration', 'num-of-doors',
          'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
          'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',
          'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',
          'horsepower', 'peak-rpm', 'city-mpg', 'highway-mpg', 'price',
          'city-L/100km', 'horsepower-binned', 'diesel', 'gas', 'diesel', 'gas'],
          dtype='object')
```

get indicator variables and assign it to data frame “dummy_variable_1”

```
[53]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
      dummy_variable_1.head()
```

```

      □
↳ -----

      KeyError                                Traceback (most recent call↳
↳ last)

      ~/conda/envs/python/lib/python3.6/site-packages/pandas/core/indexes/base.
↳ py in get_loc(self, key, method, tolerance)
      2889             try:
      -> 2890                 return self._engine.get_loc(key)
      2891             except KeyError:

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↳ _get_loc_duplicates()

```

```

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↳ _maybe_get_bool_indexer()

```

```

KeyError: 'fuel-type'

```

During handling of the above exception, another exception occurred:

```

KeyError                                Traceback (most recent call↳
↳ last)

<ipython-input-53-5c0c9c9ea224> in <module>
----> 1 dummy_variable_1 = pd.get_dummies(df["fuel-type"])
      2 dummy_variable_1.head()

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/frame.py in
↳ __getitem__(self, key)
    2973         if self.columns.nlevels > 1:
    2974             return self._getitem_multilevel(key)
-> 2975         indexer = self.columns.get_loc(key)
    2976         if is_integer(indexer):
    2977             indexer = [indexer]

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/indexes/base.
↳ py in get_loc(self, key, method, tolerance)
    2890         return self._engine.get_loc(key)
    2891         except KeyError:
-> 2892         return self._engine.get_loc(self.
↳ _maybe_cast_indexer(key))
    2893         indexer = self.get_indexer([key], method=method,↳
↳ tolerance=tolerance)
    2894         if indexer.ndim > 1 or indexer.size > 1:

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

```

```
pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.  
↳ _get_loc_duplicates()
```

```
pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.  
↳ _maybe_get_bool_indexer()
```

```
KeyError: 'fuel-type'
```

change column names for clarity

```
[48]: dummy_variable_1.rename(columns={'fuel-type-diesel': 'gas', 'fuel-type-diesel':  
↳ 'diesel'}, inplace=True)  
dummy_variable_1.head()
```

```
[48]:
```

| | diesel | gas |
|---|--------|-----|
| 0 | 0 | 1 |
| 1 | 0 | 1 |
| 2 | 0 | 1 |
| 3 | 0 | 1 |
| 4 | 0 | 1 |

We now have the value 0 to represent “gas” and 1 to represent “diesel” in the column “fuel-type”. We will now insert this column back into our original dataset.

```
[54]: # merge data frame "df" and "dummy_variable_1"  
df = pd.concat([df, dummy_variable_1], axis=1)  
  
# drop original column "fuel-type" from "df"  
df.drop("fuel-type", axis = 1, inplace=True)
```

```
↳ -----  
  
KeyError                                Traceback (most recent call↳  
↳ last)
```

```
<ipython-input-54-a92dbd6eade8> in <module>  
3  
4 # drop original column "fuel-type" from "df"  
----> 5 df.drop("fuel-type", axis = 1, inplace=True)
```

```

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/frame.py in
↳ drop(self, labels, axis, index, columns, level, inplace, errors)
    4095         level=level,
    4096         inplace=inplace,
-> 4097         errors=errors,
    4098     )
    4099

```

```

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/generic.py
↳ in drop(self, labels, axis, index, columns, level, inplace, errors)
    3913         for axis, labels in axes.items():
    3914             if labels is not None:
-> 3915                 obj = obj._drop_axis(labels, axis, level=level,
↳ errors=errors)
    3916
    3917         if inplace:

```

```

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/generic.py
↳ in _drop_axis(self, labels, axis, level, errors)
    3964         labels_missing = (axis.get_indexer_for(labels) ==
↳ -1).any()
    3965         if errors == "raise" and labels_missing:
-> 3966             raise KeyError("{} not found in axis".
↳ format(labels))
    3967
    3968         slicer = [slice(None)] * self.ndim

```

KeyError: "['fuel-type'] not found in axis"

```
[55]: df.head()
```

```

[55]:   symboling  normalized-losses      make aspiration num-of-doors  \
0         3             122  alfa-romero      std           two
1         3             122  alfa-romero      std           two
2         1             122  alfa-romero      std           two
3         2             164        audi      std           four
4         2             164        audi      std           four

   body-style drive-wheels engine-location  wheel-base  length  ...  \
0  convertible         rwd         front        88.6  0.811148  ...
1  convertible         rwd         front        88.6  0.811148  ...
2   hatchback         rwd         front        94.5  0.822681  ...

```


| | | | | | | |
|---|-------|-----|-------|------|----------|-----|
| 3 | sedan | fwd | front | 99.8 | 0.848630 | ... |
| 4 | sedan | 4wd | front | 99.4 | 0.848630 | ... |

| | highway-mpg | price | city-L/100km | horsepower-binned | diesel | gas | diesel \ |
|---|-------------|---------|--------------|-------------------|--------|-----|----------|
| 0 | 27.0 | 13495.0 | 11.190476 | Low | 0 | 1 | 0 |
| 1 | 27.0 | 16500.0 | 11.190476 | Low | 0 | 1 | 0 |
| 2 | 26.0 | 16500.0 | 12.368421 | Medium | 0 | 1 | 0 |
| 3 | 30.0 | 13950.0 | 9.791667 | Low | 0 | 1 | 0 |
| 4 | 22.0 | 17450.0 | 13.055556 | Low | 0 | 1 | 0 |

| | gas | diesel | gas |
|---|-----|--------|-----|
| 0 | 1 | 0 | 1 |
| 1 | 1 | 0 | 1 |
| 2 | 1 | 0 | 1 |
| 3 | 1 | 0 | 1 |
| 4 | 1 | 0 | 1 |

[5 rows x 33 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

Question #4:

As above, create indicator variable to the column of “aspiration”: “std” to 0, while “turbo” to 1.

```
[61]: # Write your code below and press Shift+Enter to execute
dummy_variable_2 = pd.get_dummies(df["aspiration"])
dummy_variable_2.head()

dummy_variable_2.rename(columns={'std':"aspiration-std", 'turbo':
    ↳'aspiration-turbo'}, inplace=True)

dummy_variable_2.head()
```

```
[61]: aspiration-std aspiration-turbo
0          1          0
1          1          0
2          1          0
3          1          0
4          1          0
```

[Double-click here for the solution.](#)

Question #5:

Merge the new dataframe to the original dataframe then drop the column ‘aspiration’

[62]: *# Write your code below and press Shift+Enter to execute*

```
# merge data frame "df" and "dummy_variable_2"  
df = pd.concat([df, dummy_variable_2], axis=1)  
  
# drop original column "aspiration" from "df"  
df.drop("aspiration", axis = 1, inplace=True)
```

Double-click here for the solution.

save the new csv

[]: `df.to_csv('clean_df.csv')`

Thank you for completing this notebook

<p><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.